

ASSIGNMENT - 4

1.Download the dataset

2.Import required library

In [4]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from keras.models import Model
from keras.layers import LSTM, Activation, Dense, Dropout, Input, Embedding
from keras.optimizers import RMSprop
from keras.preprocessing.text import Tokenizer
from keras.preprocessing import sequence
from keras.utils import pad_sequences
from keras.utils import to_categorical
from keras.callbacks import EarlyStopping
```

3.Read Dataset and do preprocessing

In [6]:

```
data=pd.read_csv('/content/spam.csv',encoding='latin')
```

In [7]:

```
df = pd.read_csv('/content/spam.csv',delimiter=',',encoding='latin-1')
df.head()
```

Out[7]:

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN

In [8]:

```
df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'],axis=1,inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
```

```
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    v1      5572 non-null    object
1    v2      5572 non-null    object
dtypes: object(2)
memory usage: 87.2+ KB
```

In [9]:

```
# Count of Spam and Ham values
df.groupby(['v1']).size()
```

Out[9]:

```
v1
ham      4825
spam      747
dtype: int64
```

In [10]:

```
# Label Encoding target column
X = df.v2
Y = df.v1
le = LabelEncoder()
Y = le.fit_transform(Y)
Y = Y.reshape(-1,1)
```

In [11]:

```
# Test and train split
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.15)
```

In [12]:

```
# Tokenisation function
max_words = 1000
max_len = 150
tok = Tokenizer(num_words=max_words)
tok.fit_on_texts(X_train)
sequences = tok.texts_to_sequences(X_train)
sequences_matrix = pad_sequences(sequences,maxlen=max_len)
```

4.Create Model and 5. Add Layers (LSTM, Dense-(Hidden Layers), Output)

In [13]:

```
# Creating LSTM model
inputs = Input(name='inputs', shape=[max_len])
layer = Embedding(max_words,50,input_length=max_len)(inputs)
layer = LSTM(64)(layer)
layer = Dense(256,name='FC1')(layer)
layer = Activation('relu')(layer)
layer = Dropout(0.5)(layer)
layer = Dense(1,name='out_layer')(layer)
layer = Activation('sigmoid')(layer)
model = Model(inputs=inputs,outputs=layer)
```

6.Compile the model & 7.Fit the Model

In [14]:

```
model.summary()
```

```

model.compile(loss='binary_crossentropy',optimizer=RMSprop(),metrics=['accuracy'])
model.fit(sequences_matrix,Y_train,batch_size=128,epochs=10,
        validation_split=0.2)

```

Model: "model"

Layer (type)	Output Shape	Param #
inputs (InputLayer)	[(None, 150)]	0
embedding (Embedding)	(None, 150, 50)	50000
lstm (LSTM)	(None, 64)	29440
FC1 (Dense)	(None, 256)	16640
activation (Activation)	(None, 256)	0
dropout (Dropout)	(None, 256)	0
out_layer (Dense)	(None, 1)	257
activation_1 (Activation)	(None, 1)	0

```

Total params: 96,337
Trainable params: 96,337
Non-trainable params: 0

```

```

Epoch 1/10
30/30 [=====] - 12s 311ms/step - loss: 0.3269 - accuracy: 0.8746 - val_loss: 0.1554 - val_accuracy: 0.9778
Epoch 2/10
30/30 [=====] - 12s 380ms/step - loss: 0.0819 - accuracy: 0.9794 - val_loss: 0.0461 - val_accuracy: 0.9831
Epoch 3/10
30/30 [=====] - 8s 259ms/step - loss: 0.0474 - accuracy: 0.9873 - val_loss: 0.0343 - val_accuracy: 0.9905
Epoch 4/10
30/30 [=====] - 8s 261ms/step - loss: 0.0370 - accuracy: 0.9894 - val_loss: 0.0447 - val_accuracy: 0.9895
Epoch 5/10
30/30 [=====] - 9s 310ms/step - loss: 0.0276 - accuracy: 0.9918 - val_loss: 0.0340 - val_accuracy: 0.9905
Epoch 6/10
30/30 [=====] - 8s 260ms/step - loss: 0.0224 - accuracy: 0.9929 - val_loss: 0.0402 - val_accuracy: 0.9905
Epoch 7/10
30/30 [=====] - 8s 260ms/step - loss: 0.0159 - accuracy: 0.9958 - val_loss: 0.0442 - val_accuracy: 0.9916
Epoch 8/10

```

```
30/30 [=====] - 8s 257ms/step - loss: 0.0141 - accur
acy: 0.9960 - val_loss: 0.0433 - val_accuracy: 0.9905
Epoch 9/10
30/30 [=====] - 8s 261ms/step - loss: 0.0108 - accur
acy: 0.9974 - val_loss: 0.0952 - val_accuracy: 0.9736
Epoch 10/10
30/30 [=====] - 8s 260ms/step - loss: 0.0089 - accur
acy: 0.9979 - val_loss: 0.0607 - val_accuracy: 0.9884
```

Out[14]:

```
<keras.callbacks.History at 0x7f823de6acd0>
```

8. Save the Model

In [15]:

```
model.save('sms_classifier.h5')
```

9. Test the model

In [16]:

```
test_sequences = tok.texts_to_sequences(X_test)
test_sequences_matrix = pad_sequences(test_sequences, maxlen=max_len)
```

In [17]:

```
accr = model.evaluate(test_sequences_matrix, Y_test)
27/27 [=====] - 1s 23ms/step - loss: 0.1038 - accura
cy: 0.9856
```

In [18]:

```
print('Test set\n Loss: {:.3f}\n Accuracy:
{:.3f}'.format(accr[0], accr[1]))
```

```
Test set
Loss: 0.104
Accuracy: 0.986
```